


Reliable Peer

An agent that helps you write the email you don't want to write.

“I don't want an LLM to think for me; I want it to take my thoughts and compartmentalize or augment them.”

Journal, 09.03.26

Evan Bowman
reliable-peer-ai-letter-writer.vercel.app



Start from a real conversation

Origin In class, Yuwen asked me how to email Hugh Dubberly about missing studio while seriously sick. Alex and Adin observed and took notes on how I actually answered.

The move Rather than invent a persona, we wrote a role card for behavior I had already performed — “Reliable Friend / The Know-It-All.”

Evaluation criterion “Would a friend be able to tell if it's you or an AI?” The target was my own voice, so I had ground truth to test against.

DESIGN STANCE

Augment, don't replace. A peer, not a therapist. Informational, not conversational — the agent organizes my thinking instead of substituting for it.

ROLE CARD

Team member names: Alex Li, Yuxuan Wang, Alex (Simulink), Evan, Bowman

Role Name
Reliable Friend / The Know-It-All

Purpose
To help classmates with school-related questions, specifically pertaining to relationships w/ professors and access to school resources.

Engagement Context
The friend (Alex) is seeking advice on how to write a letter to his friend about missing class due to his sickness.

Interaction Loop
1. Receive with question
2. Clarify the question
3. Provide answer / advice / information
4. End from external sources.

Behavioral Rules
- Informative
- Helpful
- Friendly
- Dismissive
- Assertive
- Assertive
- Assertive
- Assertive

Boundaries
A friend is a close friend, shaped as a respectful peer without posing the personal details.

Knowledge Base
- Knowledge of classmates (high school)
- Knowledge of context → the Bachelor's degree program.

Output
Dynamic feedback: students respond to the user's request in order of how the question is presented.

In-class role card artifact · 09.08.26

A paper artifact becomes a spec

PURPOSE

Help classmates with school questions. Informational, not conversational.

INTERACTION LOOP

4 stages: suggest → diagnose → clarify → offer alternative.

BEHAVIORAL RULES


High-level, short, single paragraph, casual tone, first names in body.

DOES NOT DO

No over-sympathizing, no fabrication, no follow-up questions, no detail dumps.

KNOWLEDGE BASE

Hugh Dubberly (values directness, conceptual clarity); MDes program; Tang Center medical policy; Jacobs Hall studio culture.



The nine-field role card ships inside the prototype as an inspectable specification, so a critic can read the rules that produced any given response.

Tuned against my own answers

V1 · BASELINE

Gemma 4 26B A4B · temperature 0 · role card as system instructions · no knowledge base

Scored against my real replies:

Q1 Rough — too many follow-ups, too much sympathy


Q2 Rough — too diagnostic, offered no backup

V2 · AFTER TUNING

Five added rules · stochastic raised to 2.0 · knowledge base attached

Rules that changed the behavior:

Lead every reply with empathy, no first-name greeting




Parameter sweeps in Agent Studio before the role card existed: same prompt, different models and temperatures, to learn which knobs actually move the voice.

WHAT THE ITERATION TAUGHT ME

More detailed prompting yields more detailed results — but rules also get over-interpreted.

Asked outright to “just write the message,” v2 still front-loaded empathy before drafting. The guardrail was working and overreaching at the same time, which is the part I could not have predicted from reading the role card alone.

How a request becomes a letter




DEPLOYMENT PIPELINE

AI Studio → GitHub (etb4311/Reliable-Peer-AI-Letter-Writer) → Vercel, free tier, with GEMINI_API_KEY set as an environment variable in the Vercel project. Verified loading in a signed-out browser.

Live and deployed

DEMO reliable-peer-ai-letter-writer.vercel.app

REPO github.com/etb4311/Reliable-Peer-AI-Letter-Writer



Left: the four-stage interaction loop with live rule-compliance indicators. Right: the drafted letter with one-click adjustments and a policy check.

Five scenario presets · four tone modes · export to mail or .txt · role card inspectable in-app

It works, and it still sounds like software

DID IT WORK?

Yes. Deployed on Vercel's free tier, loads for signed-out visitors, and behavior stayed consistent between AI Studio and production.

The role card translated cleanly into behavior, and the empathy-first rule is what made it read like a peer rather than a form letter.

WHAT BROKE

Naming conventions drifted across AI Studio, GitHub, and Vercel — cost a republish.

The model kept adding features that diluted the core concept.

Behind the GUI I could not see the agent's reasoning, only its output.

The result still feels pragmatic and impersonal next to a real friend.

WHAT I'D DO DIFFERENTLY

Make the exchange conversational in-app instead of form-driven.

Surface the agent's reasoning so the rules are visible while it works.

Run the actual test: give it to a classmate and see whether they can tell it isn't me.

Hold scope against the role card and refuse model-suggested additions.

WHO ARE WE ALLOWING TO DESIGN THE MODELS WE THINK WITH?

Building an agent in my own voice made the question concrete: every rule I wrote narrowed how a version of me was allowed to speak.